

The data verbs: `select` · `filter` · `mutate` · `arrange`

Tuesday, June 2, 2026

i Today: Tuesday June 2

Before class

- Perusall Assignment: [DC §10.1–10.6](#)

Plan for today

- Four new data verbs: `filter`, `arrange`, `select`, `mutate`
- Sorting verbs by *what they touch*: rows vs. columns
- Combining verbs with the pipe into real pipelines
- Formatting pipelines so they stay readable

Helpful links

- [Syllabus](#) · [AI policy](#) · [Schedule](#) · [R4DS Ch. 3](#) · [R4DS Ch. 4](#) · [DC Ch. 10](#) · [dplyr reference](#)

Recap from yesterday

We discussed the overarching **framework** and two data verbs:

Family	Input to Output	The verb that uses it
Reduction	variable to scalar	<code>summarise()</code>
Transformation	variable to variable	<i>...today:</i> <code>mutate()</code>
Data verb	data frame to data frame	all of them

We write reduction functions inside of `summarise()`. Today we will cover `mutate()`, which contains **transformation** functions, plus three verbs that reshape *which* rows and columns you keep.

A map of the verbs

We can organize dplyr by **what each verb operates on**:

Table 1: The dplyr toolkit.

Touches...	Verbs	What changes
Rows	<code>filter()</code> , <code>arrange()</code>	which rows, and their order
Columns	<code>select()</code> , <code>mutate()</code>	which columns, and new ones
Groups	<code>group_by()</code> , <code>summarise()</code>	(<i>yesterday</i>)

Tip

These all follow the same three rules: **first argument is a data frame, you name columns without quotes, and the output is a brand-new data frame.**

Rows: `filter()`

`filter()`: keep the rows you want

`filter()` keeps only the rows that pass a test (*DC* §10.2, *R4DS* §3.2.1):

```
penguins |>
  filter(species == "Gentoo")
```

```
# A tibble: 124 x 8
  species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
  <fct>   <fct>         <dbl>         <dbl>         <int>         <int>
1 Gentoo  Biscoe          46.1          13.2          211          4500
2 Gentoo  Biscoe           50          16.3          230          5700
3 Gentoo  Biscoe          48.7          14.1          210          4450
4 Gentoo  Biscoe           50          15.2          218          5700
5 Gentoo  Biscoe          47.6          14.5          215          5400
6 Gentoo  Biscoe          46.5          13.5          210          4550
```

```

7 Gentoo Biscoe      45.4      14.6      211      4800
8 Gentoo Biscoe      46.7      15.3      219      5200
9 Gentoo Biscoe      43.3      13.4      209      4400
10 Gentoo Biscoe      46.8      15.4      215      5150
# i 114 more rows
# i 2 more variables: sex <fct>, year <int>

```

The test is a **transformation** that returns TRUE/FALSE for each row; `filter()` keeps the TRUEs.

The comparison operators

These are the tests you'll use inside `filter()` (*DC* Table 10.9):

Table 2: Comparisons for `filter()`.

Operator	Means	Example
<code>==</code>	equal to	<code>species == "Gentoo"</code>
<code>!=</code>	not equal to	<code>sex != "male"</code>
<code>> >=</code>	greater (or equal)	<code>body_mass_g >= 5000</code>
<code>< <=</code>	less (or equal)	<code>flipper_length_mm < 190</code>
<code>%in%</code>	is one of	<code>island %in% c("Dream", "Biscoe")</code>

Warning

Use `==` to **test** equality, not `=`. A single `=` means “assign,” not “test equality”.

Multiple conditions

Separate tests with a comma (or `&`) for **AND**, a vertical bar `|` for **OR** (R4DS §3.2.1):

```

# AND: Adelies that are ALSO on Dream
penguins |>
  filter(species == "Adelie", island == "Dream")

# OR: any penguin that is Adelie OR Gentoo
penguins |>
  filter(species == "Adelie" | species == "Gentoo")

```

```
# %in% is the clean way to write a long OR
penguins |>
  filter(species %in% c("Adelie", "Gentoo"))
```

i Note

`filter(species == "Adelie" | "Gentoo")` does **not** work the way it reads. Each side of `|` must be a complete test.

Dropping NAs

Yesterday, `group_by(species, sex)` gave us an NA group because some penguins have a missing `sex`. Now we can drop them:

```
penguins |>
  filter(!is.na(sex)) |>
  group_by(species, sex) |>
  summarise(n = n(), .groups = "drop")
```

```
# A tibble: 6 x 3
  species    sex      n
  <fct>    <fct> <int>
1 Adelie  female   73
2 Adelie  male     73
3 Chinstrap female   34
4 Chinstrap male    34
5 Gentoo  female   58
6 Gentoo  male    61
```

`is.na(sex)` is TRUE for the missing ones; `!` flips it, so `!is.na(sex)` keeps the rows that **do** have a sex. This removes the NA group.

Rows: `arrange()`

`arrange()`: set the order

`arrange()` reorders rows without adding or removing any. This function gives the result in ascending order by default:

```
penguins |>
  arrange(body_mass_g)
```

```
# A tibble: 344 x 8
  species    island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
  <fct>      <fct>          <dbl>         <dbl>          <int>        <int>
1 Chinstrap Dream          46.9           16.6            192         2700
2 Adelie    Biscoe          36.5           16.6            181         2850
3 Adelie    Biscoe          36.4           17.1            184         2850
4 Adelie    Biscoe          34.5           18.1            187         2900
5 Adelie    Dream          33.1           16.1            178         2900
6 Adelie    Torgersen        38.6           17              188         2900
7 Chinstrap Dream          43.2           16.6            187         2900
8 Adelie    Biscoe          37.9           18.6            193         2925
9 Adelie    Dream          37.5           18.9            179         2975
10 Adelie   Dream          37            16.9            185         3000
# i 334 more rows
# i 2 more variables: sex <fct>, year <int>
```

You can wrap a column in `desc()` if you want descending order. If there are ties, to break them you can list several columns:

```
penguins |>
  arrange(species, desc(body_mass_g)) # by species, then heaviest first
```

Columns: `select()`

`select()`: keep the columns you want

`select()` is to columns what `filter()` is to rows:

```
penguins |>
  select(species, island, body_mass_g)
```

```
# A tibble: 344 x 3
  species island    body_mass_g
  <fct>    <fct>        <int>
1 Adelie  Torgersen     3750
2 Adelie  Torgersen     3800
```

```

3 Adelie Torgersen      3250
4 Adelie Torgersen      NA
5 Adelie Torgersen     3450
6 Adelie Torgersen     3650
7 Adelie Torgersen     3625
8 Adelie Torgersen     4675
9 Adelie Torgersen     3475
10 Adelie Torgersen     4250
# i 334 more rows

```

You can grab a range, drop with `!`, or rename as you go:

```

penguins |> select(species:bill_depth_mm) # a range of columns
penguins |> select(!year)                 # everything EXCEPT year
penguins |> select(mass = body_mass_g)    # rename while selecting

```

Note

`penguins` has only 8 columns, so `select()` feels optional here. On a 19-column table like `flights` from HW2, or a 76-column one, it's one of the first things you might want to try.

`select()` helpers

For wide tables, name columns by pattern instead of one-by-one:

```

penguins |> select(starts_with("bill")) # bill_length_mm, bill_depth_mm
penguins |> select(ends_with("_mm"))    # all the millimetre measurements
penguins |> select(where(is.numeric))   # every numeric column

```

Columns: `mutate()`

`mutate()`: add a new column

`mutate()` adds columns computed from existing ones, keeping every row:

```

penguins |>
  mutate(body_mass_kg = body_mass_g / 1000) |>
  select(species, body_mass_g, body_mass_kg)

```

```
# A tibble: 344 x 3
  species body_mass_g body_mass_kg
  <fct>      <int>      <dbl>
1 Adelie    3750        3.75
2 Adelie    3800        3.8
3 Adelie    3250        3.25
4 Adelie     NA        NA
5 Adelie    3450        3.45
6 Adelie    3650        3.65
7 Adelie    3625        3.62
8 Adelie    4675        4.68
9 Adelie    3475        3.48
10 Adelie   4250        4.25
# i 334 more rows
```

mutate() contains transformation functions

`mutate()` **always** involves a *transformation* function, just as `summarise()` **always** involves a *reduction* function.

```
# transformation: one value per row to mutate()
penguins |> mutate(bill_ratio = bill_length_mm / bill_depth_mm)

# reduction: one value total to summarise()
penguins |> summarise(avg_bill = mean(bill_length_mm, na.rm = TRUE))
```

Tip

Deciding between them is the same question from yesterday: does the calculation give **one value per row** (`mutate`) or **one value for the whole group** (`summarise`)?

A categorical column with ifelse()

`ifelse()` is a transformation that picks between two values per row, perfect inside `mutate()`:

```
penguins |>
  mutate(size = ifelse(body_mass_g > 4000, "large", "small")) |>
  select(species, body_mass_g, size)
```

```
# A tibble: 344 x 3
  species body_mass_g size
  <fct>      <int> <chr>
1 Adelie      3750 small
2 Adelie      3800 small
3 Adelie      3250 small
4 Adelie         NA <NA>
5 Adelie      3450 small
6 Adelie      3650 small
7 Adelie      3625 small
8 Adelie      4675 large
9 Adelie      3475 small
10 Adelie      4250 large
# i 334 more rows
```

Combining verbs

The pipe does the real work

One verb does one thing. For most questions, we will need several, chained together with `|>`:

```
penguins |>
  filter(!is.na(sex)) |>
  mutate(body_mass_kg = body_mass_g / 1000) |>
  select(species, sex, body_mass_kg) |>
  arrange(desc(body_mass_kg))
```

This can be read in English as: take penguins, **then** drop missing sex, **then** add mass in kg, **then** keep three columns, **then** sort heaviest first.

Tip

Because each verb returns a new data frame, the output of one is the input to the next. The verbs line up on the left so you can quickly skim to understand what's going on.

Putting all six together


We can use all six verbs to answer the question, “what’s the average mass of each species/sex, biggest first?”:


```
penguins |>
  filter(!is.na(sex)) |>
  group_by(species, sex) |>
  summarise(
    n      = n(),
    avg_mass = mean(body_mass_g, na.rm = TRUE),
    .groups = "drop"
  ) |>
  arrange(desc(avg_mass))
```

```
# A tibble: 6 x 4
  species sex      n avg_mass
  <fct>   <fct> <int>   <dbl>
1 Gentoo  male     61   5485.
2 Gentoo  female   58   4680.
3 Adelie  male     73   4043.
4 Chinstrap male    34   3939.
5 Chinstrap female   34   3527.
6 Adelie  female   73   3369.
```

And because the result is just a data frame, you can pipe it straight into a plot (switching |> to +):

```
penguins |>
  filter(!is.na(sex)) |>
  ggplot(aes(x = species, y = body_mass_g, fill = sex)) +
  geom_boxplot()
```



Formatting your pipelines

A multi-step pipeline is only readable if you format it. The rules that matter most:

- A space before `|>`, and let `|>` be the **last thing on the line**.
- For verbs with **named arguments** (`mutate`, `summarise`), put each argument on its **own line**, indented two spaces.
- Verbs without named arguments (`filter`, `select`, `arrange`) can stay on one line if they fit.
- Name new columns in `snake_case`: `body_mass_kg`, not `BodyMassKG`.

<pre># Strive for penguins > filter(!is.na(sex)) > group_by(species) > summarise(n = n(), avg = mean(body_mass_g, na.rm = TRUE))</pre>	<pre># Avoid penguins >filter(!is.na(sex)) >group_by(species) >summarize(n=n(),avg=mean(body_mass_g,na.rm=TRUE))</pre>
---	---

Tip

Don't hand-format forever: the **styler** package reformats a whole file for you. In RStudio, Cmd/Ctrl + Shift + P then type “styler”. Full guide: R4DS Ch. 4.

A helpful shortcut

Counting rows per group is so common it has its own verb, `count()`:

```
penguins |> count(species)

# is exactly the same as
penguins |>
  group_by(species) |>
  summarise(n = n())
```

Note

`count()` is just `group_by()` + `summarise(n = n())` bundled up. Add `sort = TRUE` to get the biggest groups first.

What we covered today

- **filter()**: keep rows that pass a test; `==` not `=`; combine with `,/|/%in%`; `!is.na()` to drop missing
- **arrange()**: reorder rows; **desc()** for descending; extra columns break ties
- **select()**: keep/drop/rename columns; **starts_with()** & friends for wide tables
- **mutate()**: add columns from transformations; the per-row twin of **summarise()**
- **The pipe**: chain verbs into a readable top-to-bottom story; pipe results into ggplot
- **Style**: space before `|>`, named args on their own lines, **snake_case**, let **styler** do the rest

Activity

Activity: wrangling the penguins

Roughly 25 minutes. Individually or with a neighbor.

- [Activity Canvas Link](#)

Format: a handout that walks each verb, then asks you to chain them. You'll match tasks to verbs, write **filter/arrange/select/mutate** pipelines on **penguins**, debug a `=` vs `==` error, build one multi-verb pipeline that answers a question, and restyle a deliberately unreadable pipeline.

To turn in: save your `.qmd` and rendered HTML.

Wrap-up & looking ahead

Tomorrow (Wednesday): no new reading, we put all six verbs together on a messier dataset and turn raw rows into summary tables and plots.

Upcoming Deadlines

- Thu 6/4, before class: [DC §12.1](#)
- Thu 6/4, 11:59 p.m.: [Project: Team formation](#)
- Fri 6/5, 11:59 p.m.: [Homework 3: Data Wrangling & Reshaping](#)

Reach out

Stuck? Office hours, or email me (cgc5478@psu.edu) or Xinyue (xpw5228@psu.edu).

Reference

[Syllabus](#) · [AI policy](#) · [Schedule](#) · [R4DS Ch. 3](#) · [R4DS Ch. 4](#) · [DC Ch. 10](#) · [dplyr reference](#)