

# Putting it together: from messy data to a summary table

Wednesday, June 3, 2026

**i** Today: Wednesday June 3

## Before class

- No new reading

## Plan for today

- No new verbs. Today is about **combining and sequencing** them
- The workflow: question → target table → plan → build → **check**
- Working on a genuinely messy dataset (**flights**)
- Handling missing values, small groups, wrong order

## Helpful links

- [Syllabus](#) · [AI policy](#) · [Schedule](#) · [R4DS Ch. 3](#) · [DC Ch. 7](#) · [DC Ch. 10](#) · [dplyr reference](#)

## What you already have

### Single-table toolkit:

Table 1: Six verbs, plus `count()` and the pipe.

| Touches        | Verbs                                              |
|----------------|----------------------------------------------------|
| <b>Rows</b>    | <code>filter()</code> , <code>arrange()</code>     |
| <b>Columns</b> | <code>select()</code> , <code>mutate()</code>      |
| <b>Groups</b>  | <code>group_by()</code> , <code>summarise()</code> |



### Tip

Today we will practice *choosing which verb, in what order, and checking that the answer is trustworthy and rigorous.*

## Workflow

1. **State the question** in plain English.
2. **Sketch the target table:** what does one row represent? what columns?
3. **Plan the verbs** as English sentences, writing one sentence per step.
4. **Build the pipeline** one verb at a time, looking at the output each time.
5. **Check it.** Does the count make sense? Are missing values handled? Does the result match your intuition/pass a sanity check?



### Note

Make sure not to skip Step 5. A pipeline that *runs* is not the same as a pipeline that's *correct*.

## A messier dataset

### flights

You saw `flights` on HW2 and Friday; it's also the running example in R4DS Ch. 3. It contains data on every flight out of NYC area airports in 2013:

```
glimpse(flights)
```

```
Rows: 336,776
```

```
Columns: 19
```

```
$ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2~
$ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
$ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
$ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 558, ~
$ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 600, ~
$ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2, -1~
$ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 849,~
$ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 851,~
```

```

$ arr_delay      <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7, -
1~
$ carrier        <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6", "~
$ flight         <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301, 4~
$ tailnum        <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N394~
$ origin         <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LGA",~
$ dest           <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IAD",~
$ air_time       <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149, 1~
$ distance       <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 733, ~
$ hour           <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6, 6~
$ minute         <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 0, 59, 0~
$ time_hour      <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-
01 0~

```

## Dirty data

This is a common example of what a real dataset often looks like immediately after you load it into R, before you've done any data cleaning:

- **336,776 rows, 19 columns:** `select()` and `n()` are much more relevant here compared to a smaller dataset like `penguins`.
- **Meaningful missing values:** cancelled flights have NA for `dep_delay` and `arr_delay`.
- **Codes, rather than long-form labels:** here, `carrier` is "UA", not "United". (Turning codes into names requires a *join*, which is later, so today we'll stick to the codes.)

```

flights |>
  summarise(
    n          = n(),
    missing_dep = sum(is.na(dep_delay))
  )

```

```

# A tibble: 1 x 2
      n missing_dep
  <int>      <int>
1 336776      8255

```

### Warning

The above result tells us that there are 8,000 rows with a missing departure delay. Every average we compute has to decide what to do with them, or it silently returns NA.

## Worked example 1

**Q: which months had the worst delays?**

**Plan:**

- We want one row per **month**, with its average departure delay and a count.
- *Group by month, then summarise the mean and n(), then arrange worst-first.*

**Target table:** month | avg\_delay | n, which should have twelve rows.

```
flights |>
  group_by(month) |>
  summarise(
    avg_delay = mean(dep_delay, na.rm = TRUE),
    n         = n(),
    .groups   = "drop"
  ) |>
  arrange(desc(avg_delay))
```

```
# A tibble: 12 x 3
  month avg_delay    n
  <int>   <dbl> <int>
1     7    21.7 29425
2     6    20.8 28243
3    12    16.6 28135
4     4    13.9 28330
5     3    13.2 28834
6     5    13.0 28796
7     8    12.6 29327
8     2    10.8 24951
9     1    10.0 27004
10    9     6.72 27574
11   10     6.24 28889
12   11     5.44 27268
```

## What just happened

We used **three** verbs in sequence, and two safety habits from this week:

- `na.rm = TRUE` so the cancelled flights don't give NAs for the means.

- `n = n()` so we can see each average rests on around 25,000 flights at least, rather than a small number.

#### Tip

Summer months (6, 7) and December rise to the top, which matches what you'd expect from thunderstorms and holiday travel. This is a good place to think about common sense/sanity checking; if January had come out worst by a mile (no thunderstorms and little holiday travel), it might suggest something is wrong with your code. At the very least, it's a signal to double check your code/investigate further.

## Worked example 2

### Q: worst JFK summer destinations?

**Question:** for flights leaving **JFK in summer**, which **destination** had the worst average arrival delay, among destinations with **enough flights to trust the result**? Let's set an arbitrary cutoff at 100 flights for simplicity.

#### Plan:

- *Filter* to JFK, summer months.
- *Group by* destination, *summarise* mean arrival delay + `n()`.
- *Filter again* to drop tiny destinations.
- *Arrange* worst-first.

#### Execution

```
flights |>
  filter(origin == "JFK", month %in% c(6, 7, 8)) |>
  group_by(dest) |>
  summarise(
    avg_arr_delay = mean(arr_delay, na.rm = TRUE),
    n              = n(),
    .groups        = "drop" # <-- removes grouping metadata from your data frame
  ) |>
  filter(n >= 100) |>
  arrange(desc(avg_arr_delay))
```

```
# A tibble: 49 x 3
  dest   avg_arr_delay     n
  <chr>         <dbl> <int>
1 CMH           44.4   194
2 CVG           39.4   267
3 IND           30.8   144
4 PDX           29.6   265
5 BNA           28.8   184
6 DEN           28.6   184
7 CLE           28.2   199
8 ORF           26.7   128
9 PBI           25.9   368
10 JAX           22.1   365
# i 39 more rows
```

#### Note

Notice that `filter()` appears **twice**, operating once on the raw rows (which flights), and once on the summary table (which destinations are “big enough”).

### The necessity of the second filter (`.scrollable .nostretch`)

Without `filter(n >= 100)`, the “worst destination” could be whichever one had a single very-late flight (isn’t necessarily a problem with this dataset, but we got lucky). **A mean from a group of 3 is not very reliable.**

- If we gained more samples, the answer is much more likely to shift significantly in the case of a small sample size.

It’s a good idea to have verification steps built-in to all of your pipelines when possible; always including the count column is how you *catch* the problem before it ends up in a report that you have to retract later.

## Worked example 3

### Q: do long hauls leave later?

Now we can leverage `mutate()` before piping our transformed data into a plot.

**English plan:** *mutate* a long/short label from `distance`, *group by* it, *summarise* the average delay, *then plot* the summary table.

```

flights |>
  mutate(haul = ifelse(distance > 1000, "long", "short")) |>
  group_by(haul) |>
  summarise(
    avg_delay = mean(dep_delay, na.rm = TRUE),
    n         = n(),
    .groups   = "drop"
  )

```

```

# A tibble: 2 x 3
  haul avg_delay     n
  <chr>   <dbl> <int>
1 long      11.5 147105
2 short     13.5 189671

```


## From summary table to plot

Because the summary is just a (tiny) data frame, it's already glyph-ready, so we can pipe it into `ggplot()`:

```

flights |>
  mutate(haul = ifelse(distance > 1000, "long", "short")) |>
  group_by(haul) |>
  summarise(avg_delay = mean(dep_delay, na.rm = TRUE), .groups = "drop") |>
  ggplot(aes(x = haul, y = avg_delay)) +
  geom_col()

```



## Getting it right

### Order matters

The same verbs in a different order answer different questions:

#### Filter, *then* summarise

“Average delay *of JFK flights*.” (This is more computationally efficient)

```
flights |>
  filter(origin == "JFK") |>
  summarise(
    avg = mean(dep_delay,
               na.rm = TRUE)
  )
```

#### Summarise, *then* filter

“Of the per-origin averages, keep JFK’s.”

```
flights |>
  group_by(origin) |>
  summarise(
    avg = mean(dep_delay,
```

```

      na.rm = TRUE)
) |>
filter(origin == "JFK")

```

Both are valid, and both give the same number in this case. However, they won't always be. The trick is knowing **which question you're actually asking** before you write it.

## A case when the output changes with order

```

flights |>
  filter(dep_delay > 60) |>
  summarise(avg_delay = mean(arr_delay, na.rm = TRUE))

```

```

# A tibble: 1 x 1
  avg_delay
  <dbl>
1      119.

```

```

flights |>
  summarise(avg_delay = mean(arr_delay, na.rm = TRUE)) |>
  filter(avg_delay > 60)

```

```

# A tibble: 0 x 1
# i 1 variable: avg_delay <dbl>

```

## The three checks

Before trusting any wrangled result, run this checklist:

Table 2: Your verification checklist.

| Check                                | The verb/habit                             |
|--------------------------------------|--------------------------------------------|
| Did missing values get handled?      | <code>na.rm</code> , <code>!is.na()</code> |
| Is each number based on enough rows? | <code>n = n()</code>                       |
| Does the answer pass a smell test?   | compare to what you know                   |

 Warning

A pipeline that runs without error has cleared the *lowest* bar; that doesn't *necessarily* mean everything is fine.

## What we did today

- **No new verbs:** just the six from Monday and Tuesday, sequenced
- **The workflow:** question → target table → English plan → build incrementally → check
- **Real mess:** missing delays, code-not-name columns, tiny groups
- **filter() in two roles:** subsetting raw rows *and* trimming a summary table
- **Order matters:** filter-then-summarise ≠ summarise-then-filter
- **The full arc:** raw rows → glyph-ready summary → plot

## Activity

### Activity: a flights investigation

Roughly 25 minutes. Individually or with a neighbor.

- [Activity Canvas Link](#)

**Format:** you'll answer a series of questions about **flights** by building pipelines, but each one starts with writing the **plan in words first**. You'll compare two pipelines that look alike but answer different questions, catch a silently-dropped NA, and finish with one open-ended question where you choose the verbs and **report your sanity checks**.

**To turn in:** save your .qmd and rendered HTML.

## Wrap-up & looking ahead

**Tomorrow (Thursday):** new reading and new content, reshaping data between **wide and narrow** with `pivot_longer()` and `pivot_wider()`.

### Upcoming Deadlines

- Thu 6/4, before class: [DC §12.1](#)
- Thu 6/4, 11:59 p.m.: [Project: Team formation](#)
- Fri 6/5, 11:59 p.m.: [Homework 3: Data Wrangling & Reshaping](#)

**Reach out**

Stuck? Office hours, or email me ([cgc5478@psu.edu](mailto:cgc5478@psu.edu)) or Xinyue ([xpw5228@psu.edu](mailto:xpw5228@psu.edu)).

**Reference**

[Syllabus](#) · [AI policy](#) · [Schedule](#) · [R4DS Ch. 3](#) · [DC Ch. 10](#) · [dplyr reference](#)