

Iteration: `map()`, `walk()`, and saving many outputs

Wednesday, June 17, 2026

i Today: Wednesday June 17

Before class

- **Reading Quiz** (Canvas): [R4DS Ch. 26](#) §26.4–26.5 — Iteration

Plan for today

- From `across()` (columns) to `map()` (any list or vector)
- The `map()` family and typed variants (`map_dbl()`, `map_chr()`, ...)
- `walk()`: when you only want the side effect
- **List-columns** with `group_nest()`; building paths with `str_glue()`
- Two inputs at once: `map2()` / `walk2()`; saving many CSVs and plots
- In-class activity

Helpful links

- [Syllabus](#) · [AI policy](#) · [Schedule](#) · [R4DS Ch. 26](#) · [purrr reference](#)

The `map()` family

From `across()` to `map()`

Last class, `across()` repeated an action over the **columns** of a data frame.

`map()` is more general: it applies a function to **every element of any list or vector** and returns the results.

```
map(1:3, \(x) x^2)
```

```
[[1]]  
[1] 1
```

```
[[2]]  
[1] 4
```

```
[[3]]  
[1] 9
```

`map(x, f)` calls `f` once per element of `x` and collects the answers into a **list**. (Same `\(x)` anonymous-function shorthand as last class.)

Typed variants return vectors

A list isn't always what you want. The typed variants return an atomic vector, and error if the result isn't that type:

```
map_dbl(1:3, \(x) x^2)          # double vector
```

```
[1] 1 4 9
```

```
map_chr(c("a", "bb"), \(x) str_c(x, "!"))  # character vector
```

```
[1] "a!" "bb!"
```

```
map_int(list(1:5, 1:10), length)          # integer vector
```

```
[1] 5 10
```

Function	Returns
<code>map()</code>	a list
<code>map_dbl()</code> / <code>map_int()</code>	numeric / integer vector
<code>map_chr()</code>	character vector
<code>map_lgl()</code>	logical vector

A data frame is a list of columns

Because a data frame is really a list of columns, `map()` walks over its columns:

```
nums <- diamonds |> select(carat, depth, price)
map_dbl(nums, median)
```

```
carat depth price
0.7    61.8 2401.0
```

So `map_dbl(df, median)` is another route to the same place as `summarize(across(everything(), median))`. Different tools, same idea: *do this to each thing*.

Note

You may see older code with `~ .x + 1` instead of `\(x) x + 1`. That tilde syntax only works inside tidyverse functions and always names the argument `.x`. Prefer the base `\(x)` form.

`map()` reads many files

A classic use: read a folder of files and stack them. Get the paths, `map()` `read_csv` over them, then `list_rbind()` to bind the list of data frames into one:

```
paths <- list.files("data/gapminder", pattern = "\\*.csv$", full.names = TRUE)

gapminder <- paths |>
  map(read_csv) |> # a list of data frames, one per file
  list_rbind()     # stacked into a single data frame
```

This class we tackle the **opposite** problem: taking R objects and **saving them to many files**.

Saving many outputs

`walk()`: `map()` that throws the output away

When you call a function only for its **side effect** (writing a file, printing, saving a plot), you don't care about the return value. `walk()` is `map()` without the output:

```
walk(c("first", "second", "third"), \(x) cat(x, "\n"))
```

```
first
second
third
```

Compare: `map()` would also return a list of NULLs, cluttering your console. `walk()` returns its input invisibly.

Example: many files into a database

When data won't all fit in memory, load it into a database, then query the bits you need. The pattern: a **template** table, then **append** each file.

```
con <- DBI::dbConnect(duckdb::duckdb())

# 1. build an empty table with the right columns/types from one file
template <- readxl::read_excel(paths[[1]])
template$year <- 1952
DBI::dbCreateTable(con, "gapminder", template)

# 2. a function that reads ONE file and appends it
append_file <- function(path) {
  df <- readxl::read_excel(path)
  df$year <- parse_number(basename(path))
  DBI::dbAppendTable(con, "gapminder", df)
}

# 3. walk() it over every path (we don't need a return value)
paths |> walk(append_file)
```

`walk()` here, not `map()`, because we only care about the rows landing in the database.

List-columns with `group_nest()`

To save one file per group, first split the data into per-group tibbles. `group_nest()` packs each group's rows into a **list-column** called `data`:

```
by_clarity <- diamonds |> group_nest(clarity)
by_clarity
```

```
# A tibble: 8 x 2
  clarity      data
  <ord>    <list<tibble[,9]>>
1 I1      [741 x 9]
2 SI2     [9,194 x 9]
3 SI1     [13,065 x 9]
4 VS2     [12,258 x 9]
5 VS1     [8,171 x 9]
6 VVS2    [5,066 x 9]
7 VVS1    [3,655 x 9]
8 IF      [1,790 x 9]
```

Each element of `data` is a full tibble:

```
by_clarity$data[[1]]
```

```
# A tibble: 741 x 9
  carat cut      color depth table price      x      y      z
  <dbl> <ord>    <ord> <dbl> <dbl> <int> <dbl> <dbl> <dbl>
1  0.32 Premium  E      60.9   58   345  4.38  4.42  2.68
2  1.17 Very Good J      60.2   61  2774  6.83  6.9   4.13
3  1.01 Premium  F      61.8   60  2781  6.39  6.36  3.94
4  1.01 Fair     E      64.5   58  2788  6.29  6.21  4.03
5  0.96 Ideal    F      60.7   55  2801  6.37  6.41  3.88
6  1.04 Premium  G      62.2   58  2801  6.46  6.41  4
7  1 Fair       G      66.4   59  2808  6.16  6.09  4.07
8  1.2 Fair     F      64.6   56  2809  6.73  6.66  4.33
9  0.43 Very Good E      58.4   62   555  4.94  5     2.9
10 1.02 Premium  G      60.3   58  2815  6.55  6.5   3.94
# i 731 more rows
```

Build the output paths with `str_glue()`

Add a column of file names. `str_glue()` interpolates the value of `clarity` into each string:

```
by_clarity <- by_clarity |>
  mutate(path = str_glue("diamonds-{clarity}.csv"))
by_clarity
```

```
# A tibble: 8 x 3
  clarity      data path
  <ord>    <list<tibble[,9]>> <glue>
1 I1      [741 x 9] diamonds-I1.csv
2 SI2     [9,194 x 9] diamonds-SI2.csv
3 SI1     [13,065 x 9] diamonds-SI1.csv
4 VS2     [12,258 x 9] diamonds-VS2.csv
5 VS1     [8,171 x 9] diamonds-VS1.csv
6 VVS2     [5,066 x 9] diamonds-VVS2.csv
7 VVS1     [3,655 x 9] diamonds-VVS1.csv
8 IF      [1,790 x 9] diamonds-IF.csv
```

We now have everything we need side by side: the **data** to write and the **path** to write it to.

Two inputs at once: `map2()` / `walk2()`

Saving needs **two** things that vary together — the data *and* the path. By hand that's:

```
write_csv(by_clarity$data[[1]], by_clarity$path[[1]])
write_csv(by_clarity$data[[2]], by_clarity$path[[2]])
# ... eight times
```

`map2()` varies the **first two** arguments together; `walk2()` does that and discards the output:


```
walk2(by_clarity$data, by_clarity$path, write_csv)
```

One line writes all eight CSVs. `walk2(x, y, f)` calls `f(x[[i]], y[[i]])` for each `i`.

Saving many plots

Same recipe. First a function that builds the plot for one data frame:

```
carat_histogram <- function(df) {
  ggplot(df, aes(x = carat)) + geom_histogram(binwidth = 0.1)
}
carat_histogram(by_clarity$data[[1]])
```



Saving many plots

Use `map()` to make a **list-column of plots**, and `str_glue()` for the paths:

```
by_clarity <- by_clarity |>
  mutate(
    plot = map(data, carat_histogram),
    path = str_glue("clarity-{clarity}.png")
  )
```

Then `walk2()` + `ggsave()` writes them all:

```
walk2(
  by_clarity$path,
  by_clarity$plot,
  \(path, plot) ggsave(path, plot, width = 6, height = 6)
)
```

Reading files, writing files, drawing plots — all the same `map/walk` shape.

Why this matters

Iteration is a superpower

The whole arc of this chapter is one idea applied three ways:

Task	Tool
do something to each column	<code>across()</code>
do something to each element	<code>map()</code> / <code>map_dbl()</code> / ...
do something for its side effect	<code>walk()</code> / <code>walk2()</code>

“If you know the right iteration technique, you can easily go from fixing one problem to fixing all the problems.” — R4DS

What about for loops?

You may have noticed we haven’t written a single `for` loop.

R’s data-analysis orientation changes how we iterate: most of the time there’s already an idiom — do something to each **column** (`across`), each **group** (`group_by`), or each **element** (`map`).

`for` loops *do* exist and you’ll meet them in **wild-caught code** (and a bit of base R later). But reach for a functional tool first: it’s shorter, harder to get wrong, and says *what* you mean, not just *how*.

Activity

In-class activity

Open `day22-activity.R` in RStudio and work through the parts in order.

- **Part 1:** `map()` and its typed variants; anonymous functions; `map()` over the columns of a data frame.
- **Part 2:** `group_nest()` to build a list-column, `str_glue()` for paths, and `walk2()` to write one CSV per group (to a temp folder).
- **Part 3 (challenge):** write a plot function, build a list-column of plots with `map()`, and save them all with `walk2()` + `ggsave()`.

[Activity file](#)

Wrap-up & looking ahead

What we covered

- `map(x, f)` applies `f` to each element of `x` → a list
- Typed variants (`map_dbl`, `map_chr`, `map_int`, `map_lgl`) → vectors
- `walk()` = `map()` for side effects (writing, saving)

- `group_nest()` makes a per-group **list-column**
- `str_glue()` builds output paths from data
- `map2()` / `walk2()` vary **two** inputs together (data + path)
- Reading files, writing CSVs, saving plots — all the same shape

Upcoming

- **Thu Jun 18:** strings, dates & intro to regex + **Project oral check-in #1**
- **Reading before Thu:** *Data Computing* §17.1–17.2

Reach out

Stuck? Office hours, or email me (cgc5478@psu.edu) or Xinyue (xpw5228@psu.edu).

Reference

[Syllabus](#) · [AI policy](#) · [Schedule](#) · [R4DS Ch. 26](#) · [purrr reference](#)